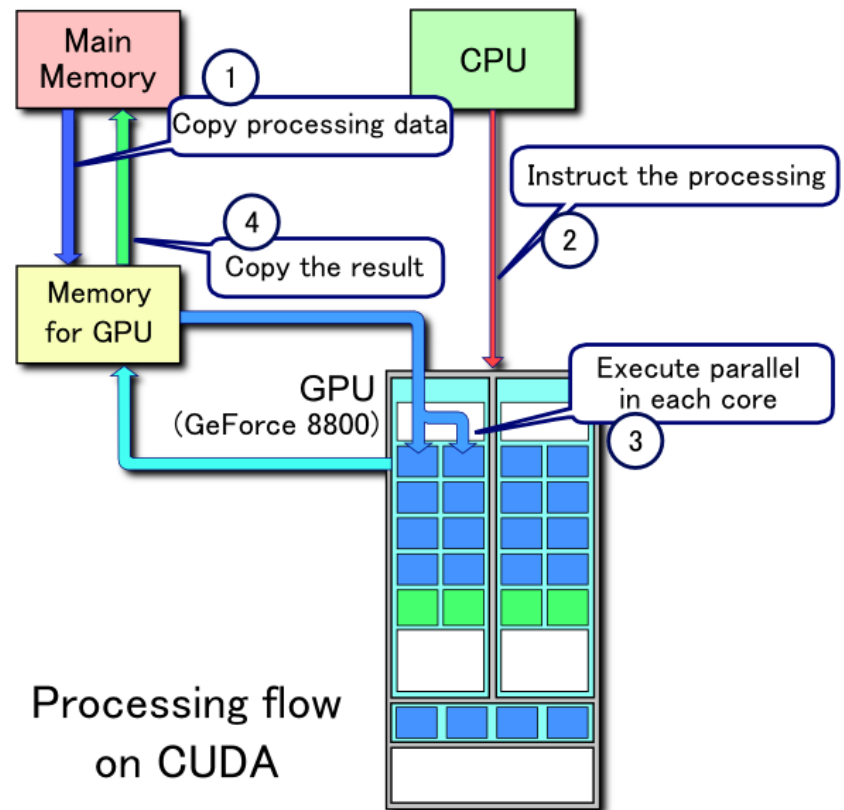


Dynamic Parallelism

Compute Unified Device Architecture

- Hybrid CPU/GPU Code
- Low latency code is running on CPU
 - Result immediately available
- High latency, high throughput code is running on GPU
 - Result on bus
 - GPU has many more cores than CPU



Types of Parallelism

- Different Types of Parallelism
 - Task parallelism (Task farming, Divide and conquer)
 - Problem is divided to tasks, which are processed independently.
 - Data parallelism (SPMD)
 - Same operation is performed over many data items
 - Other types of parallelism
 - Event driven, ...


GPU Execution Model

- Parallelism in GPU
 - Data parallelism
 - The same kernel is executed by many threads
 - Thread process one data item
 - Limited task parallelism
 - Multiple kernels executed simultaneously (since Fermi)
 - But we do not have
 - Any means of kernel-wide synchronization (barrier)
 - Any guarantees that two blocks/kernels will actually run concurrently

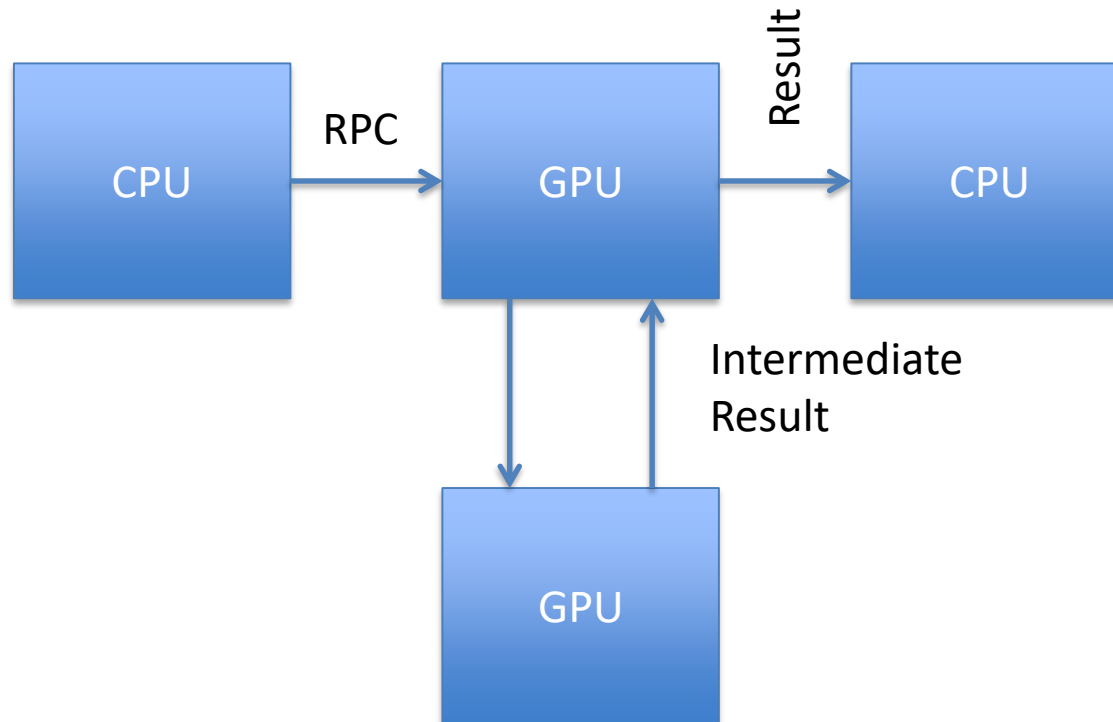
What is Dynamic Parallelism

- The ability to launch new kernels from the GPU
 - Dynamically
 - based on run-time data
 - Simultaneously
 - from multiple threads at once
 - Independently
 - each thread can launch a different grid
- Introduced with CUDA 5.0 and compute capability 3.5 and up

Execution Model (Overview)



Fermi: Only CPU can generate GPU work






Kepler: GPU can generate work for itself

Dynamic Parallelism

- Allows program flow to be controlled by GPU
- Allows recursion and subdivision of problems
- Interesting when data is not uniformly distributed
- Dynamic parallelism can launch additional threads in interesting areas
- Allows higher resolution in critical areas without slowing down others

Problematic Cases

- Unsuitable Problems for GPUs
 - Processing irregular data structures
 - Trees, graphs, ...
 - Regular structures with irregular processing workload
 - Difficult simulations, iterative approximations
 - Iterative tasks with explicit synchronization

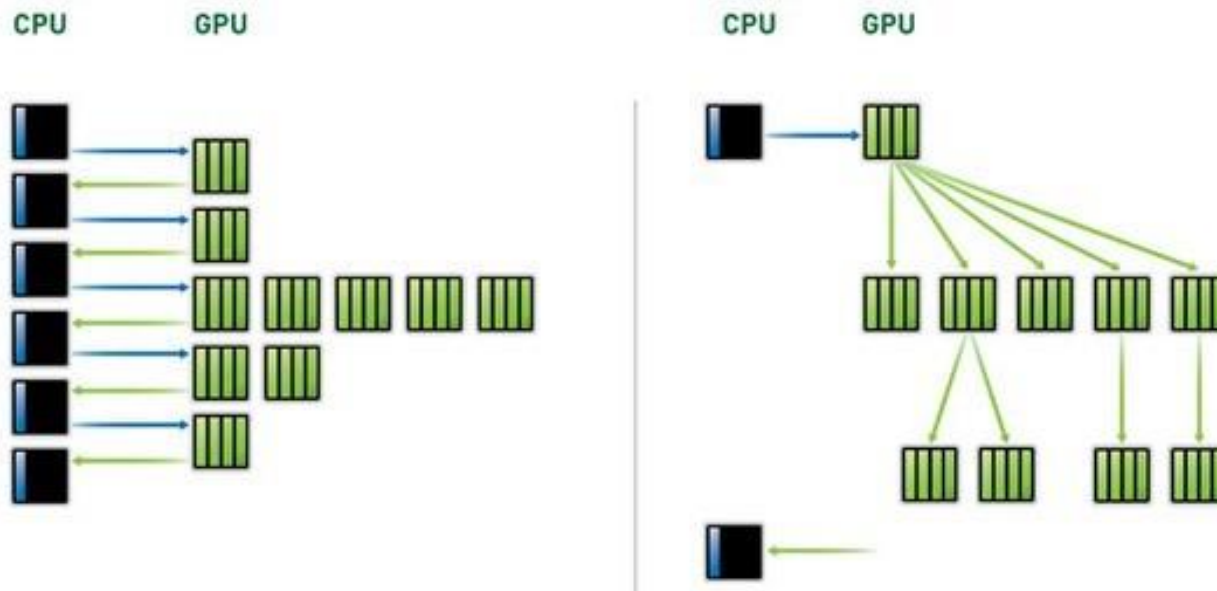


Problematic Cases

- Solutions
 - Iterative kernel execution
 - Usually applicable only for cases when there are none or few dependencies between the subsequent kernels
 - The state (or most of it) is kept on the GPU
 - Mapping irregular structures to regular grids
 - May be too fine/coarse grained
 - Not always possible
 - 2-phase Algorithms
 - First phase determines the amount of work (items, ...)
 - Second phase process tasks mapped by first phase


Dynamic Parallelism

- Dynamic Parallelism Purpose
 - The device does not need to synchronize with host to issue new work to the device
 - Irregular parallelism may be expressed more easily

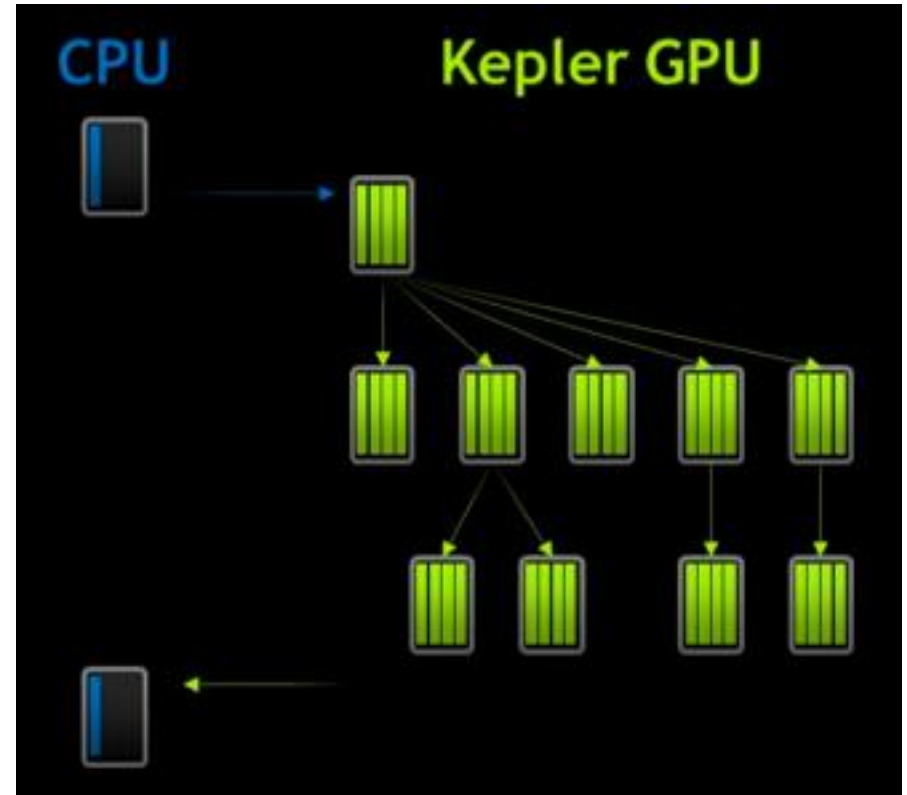


Dynamic Parallelism

- Without Dynamic Parallelism
 - Data travels back and forth between the CPU and GPU many times.
 - This is because of the inability of the GPU to create more work on itself depending on the data.



- With Dynamic Parallelism:
 - GPU can generate work on itself based on intermediate results, without involvement of CPU.
 - Permits Dynamic Run Time decisions.
 - Leaves the CPU free to do other work, conserves power.



Dynamic Parallelism

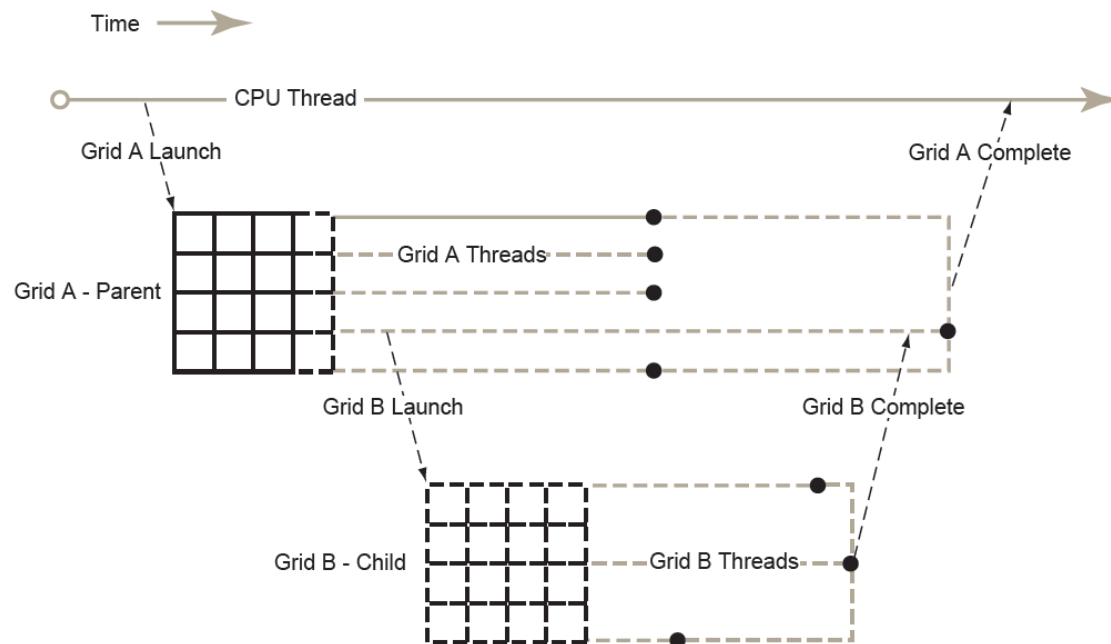
- How It Works
 - Portions of CUDA runtime are ported to device side
 - Kernel execution
 - Device synchronization
 - Streams, events, and async memory operations
 - Kernel launches are asynchronous
 - No guarantee the child kernel starts immediately
 - Synchronization points may cause context switch
 - Entire blocks are switched on a SMP

Dynamic Parallelism


- CUDA Dynamic Parallelism
 - New feature presented in CC 3.5 (Kepler)

— (

|




Dynamic Work Generation




Initial Grid

*Statically assign conservative
worst-case grid*

Fixed Grid



*Dynamically assign performance
where accuracy is required*



Dynamic Grid

Dynamic Parallelism

- Example

```
__global__ void child_launch(int *data) {  
    data[threadIdx.x] = data[threadIdx.x]+1;  
}
```

```
__global__ void parent_launch(int *data) {  
    data[threadIdx.x] = threadIdx.x;  
    __syncthreads();  
    if (threadIdx.x == 0) {  
        child_launch<<< 1, 256 >>>(data);  
        cudaDeviceSynchronize();  
    }  
    __syncthreads();  
}
```

```
void host_launch(int *data) {  
    parent_launch<<< 1, 256 >>>(data);  
}
```

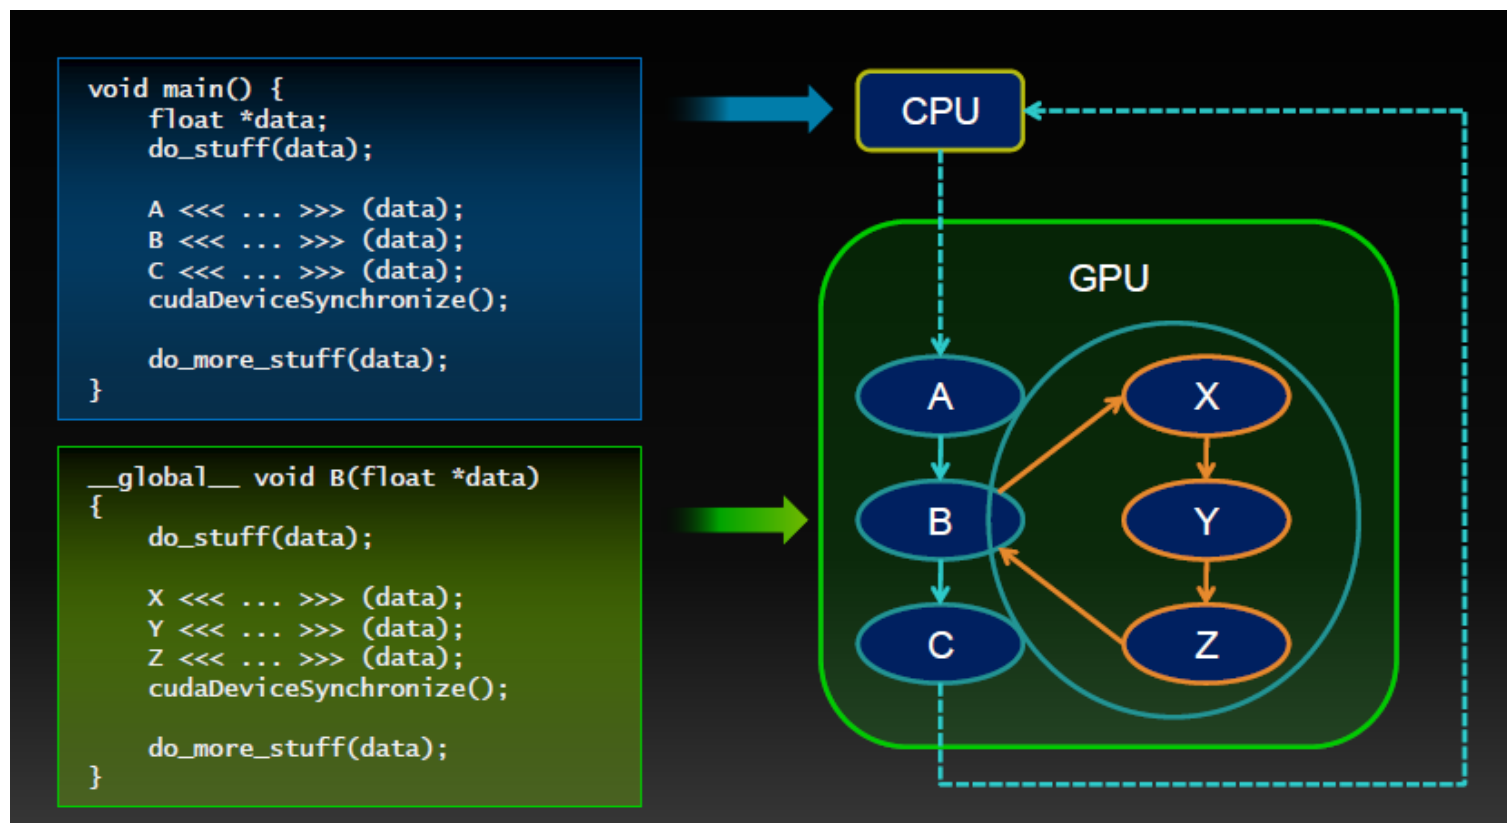
Thread 0 invokes a grid of child threads

Synchronization does not have to be invoked by all threads

Device synchronization does not synchronize threads in the block

Dynamic Parallelism

- Nested Dependencies



Source: NVIDIA

EASY TO PARALLELIZE PROGRAM

```
for i = 1 to N
  for j = 1 to M
    do_something(i, j)
  next j
next i
```

EASY TO PARALLELIZE PROGRAM

Serial Code

```
for i = 1 to N
  for j = 1 to M
    do_something(i, j)
  next j
next i
```

Parallel Code

```
__global__ void callKernel(...)
{
  do_something (blockIdx.x,
                threadIdx.x)
}
```

```
void main()
{
  callKernel<<< N, M >>>(...);
}
```

threadIdx.x

threadIdx.x

threadIdx.x

threadIdx.x



blockIdx.x = 0

blockIdx.x = 1

blockIdx.x = 2

blockIdx.x = 3

EASY TO PARALLELIZE PROGRAM


Serial Code

```
for i = 1 to N
  for j = 1 to M
    do_something(i, j)
  next j
next i
```

Parallel Code

```
__global__ void callKernel( ... )
{
  do_something (blockIdx.y,
               blockIdx.x)
}
```

```
void main()
{
  dim3 grid(N, M)
  callKernel<<< grid, 1 >>>( ... );
}
```



EASY TO PARALLELIZE PROGRAM


Serial Code

```
for i = 1 to N
  for j = 1 to M
    do_something(i, j)
  next j
next i
```

Parallel Code

```
__global__ void callKernel( ... )
{
  do_something (threadIdx.y,
               threadIdx.x)
}
```

```
void main()
{
  dim3 block(N, M)
  callKernel<<< 1, block >>>(...);
}
```



EASY TO PARALLELIZE PROGRAM

Serial Code

```
for i = 1 to N
  for j = 1 to M
    do_something(i, j)
  next j
next i
```

Device

Grid 1

Block (0, 0)	Block (1, 0)	Block (2, 0)
Block (0, 1)	Block (1, 1)	Block (2, 1)

Block (1, 1)

Thread (0, 0)	Thread (1, 0)	Thread (2, 0)	Thread (3, 0)	Thread (4, 0)
Thread (0, 1)	Thread (1, 1)	Thread (2, 1)	Thread (3, 1)	Thread (4, 1)
Thread (0, 2)	Thread (1, 2)	Thread (2, 2)	Thread (3, 2)	Thread (4, 2)

Parallel Code

```
__global__ void callKernel( ... )
{
  int i = ...;   int j = .....;
  do_something (i, j)
}
```

```
void main()
{ dim3 grid(A, B)
  dim3 block(X = N/A, Y= M/B)
  callKernel<<< grid, block >>>(...);
}
```

DIFFICULT TO PARALLELIZE PROGRAM


Serial Code

```
for i = 0 to N
  for j = 0 to x[i]
    do_something(i, j)
  next j
next i
```

Bad alternative : Idle Threads

```
__global__ void callKernel(int x[], ...)
{
  int i = blockIdx.x;
  int j = threadIdx.x;
  if ( j < x[i] )
    do_something (i, j)
}
```


```
void main()
{
  callKernel<<< N, max(x) >>>(x...);
}
```



DIFFICULT TO PARALLELIZE PROGRAM

Serial Code

```
for i = 1 to N
  for j = 1 to x[i]
    do_something(i, j)
  next j
next i
```



Dynamic Parallel Code

```
__global__ void ChildKernel(int i) {
  int j = threadIdx.x;
  do_something(i, j);
}
```

```
__global__ void kernel( ... )
{
  int i = threadIdx.x;
  Childkernel<<< 1, x[i]>>>(i)
}
```

```
void main()
{
  kernel<<< 1, N >>>(x);
}
```


Dynamic Parallelism

- Nested Dependencies -
`cudaDeviceSynchronize ()`
- Can be used inside a kernel
- Synchronizes all launches by any kernel in block
- Does NOT imply `__syncthreads()`!

Dynamic Parallelism

- Kernel launch implies memory sync operation
- Child sees state at time of launch
- Parent sees child writes after sync
- Local and shared memory are private, cannot be shared with children