

CUDA Programming

Hello Program

Outline

☐ CUDA Programming

- ☐ Functions Qualifiers

- ☐ Built-in Device Variables

- ☐ Variable Qualifiers

☐ Addition on the device

- ☐ Moving to parallel using blocks

- ☐ Moving to parallel using threads

- ☐ Combining blocks and threads

Cuda Programming

- **Kernels** are C functions with some restrictions

- Can only access GPU memory
- Must have void return type
- No variable number of arguments ("varargs")
- Not recursive
- No static variables

local variables / متغيرات محلية
تبقى حتى انتهاء
العملية (لا تُحذف)
حتى بعد الانتهاء

- Function arguments automatically copied from CPU to GPU memory

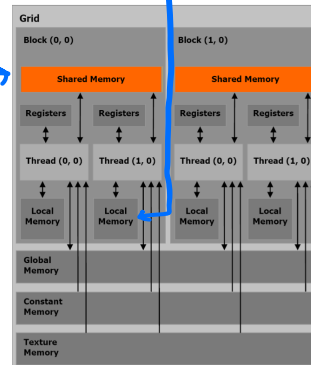
Function Qualifiers

- **__global__** : invoked from within host (CPU) code,
 - cannot be called from device (GPU) code
 - must return **void**
- **__device__** : called from other GPU functions,
 - cannot be called from host (CPU) code
- **__host__** : can only be executed by CPU, called from host
- **__host__** and **__device__** qualifiers can be combined
 - Sample use: overloading operators
 - Compiler will generate both CPU and GPU code

Variable Qualifiers (GPU code)

↳ ex: `_shared_ int x;`

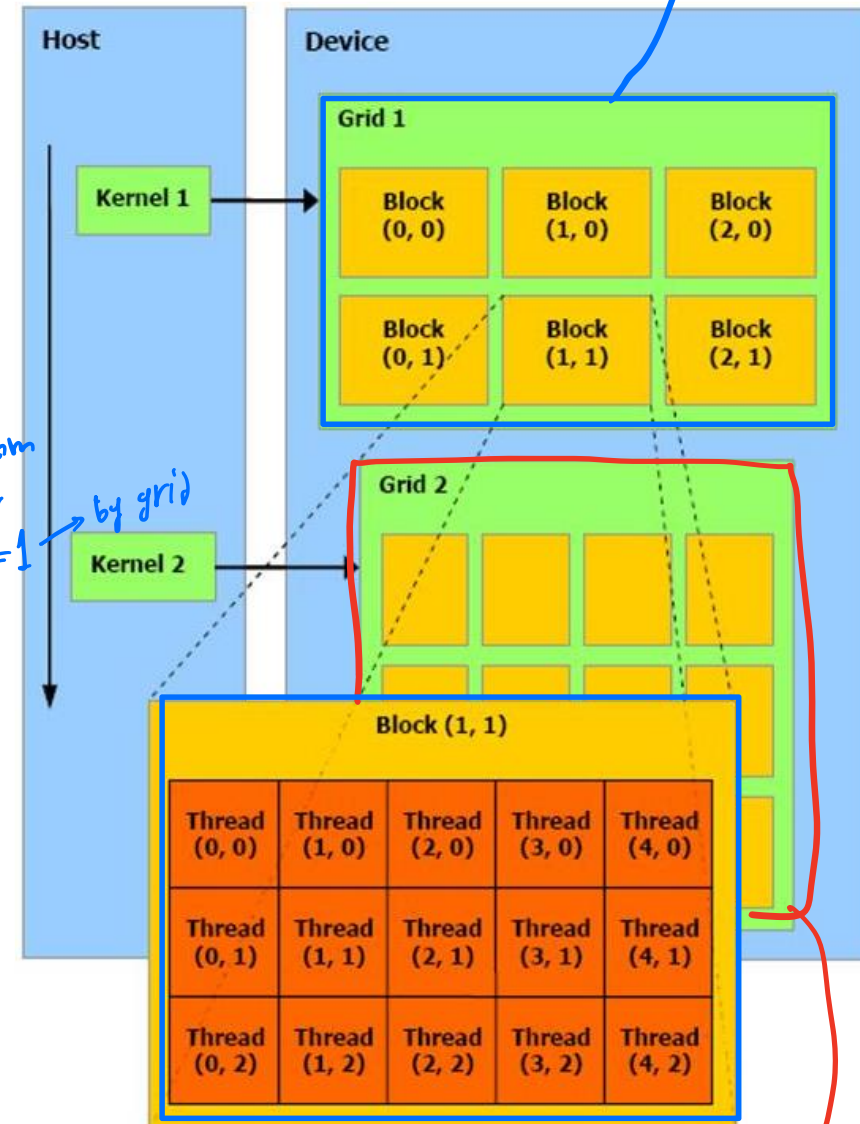
- **`__device__`**
 - Stored in device memory (large, high latency, no cache)
 - Allocated with `cudaMalloc` (`__device__` qualifier implied)
 - Accessible by all threads
 - Lifetime: application
- **`__shared__`**
 - Stored in on-chip shared memory (very low latency)
 - Allocated by execution configuration or at compile time
 - Accessible by all threads in the same thread block
 - Lifetime: kernel execution
- Unqualified variables:
 - Scalars and built-in vector types are stored in registers
 - Arrays of more than 4 elements stored in device memory



CUDA Built-in Device Variables

All `__global__` and `__device__` functions have access to these automatically defined variables

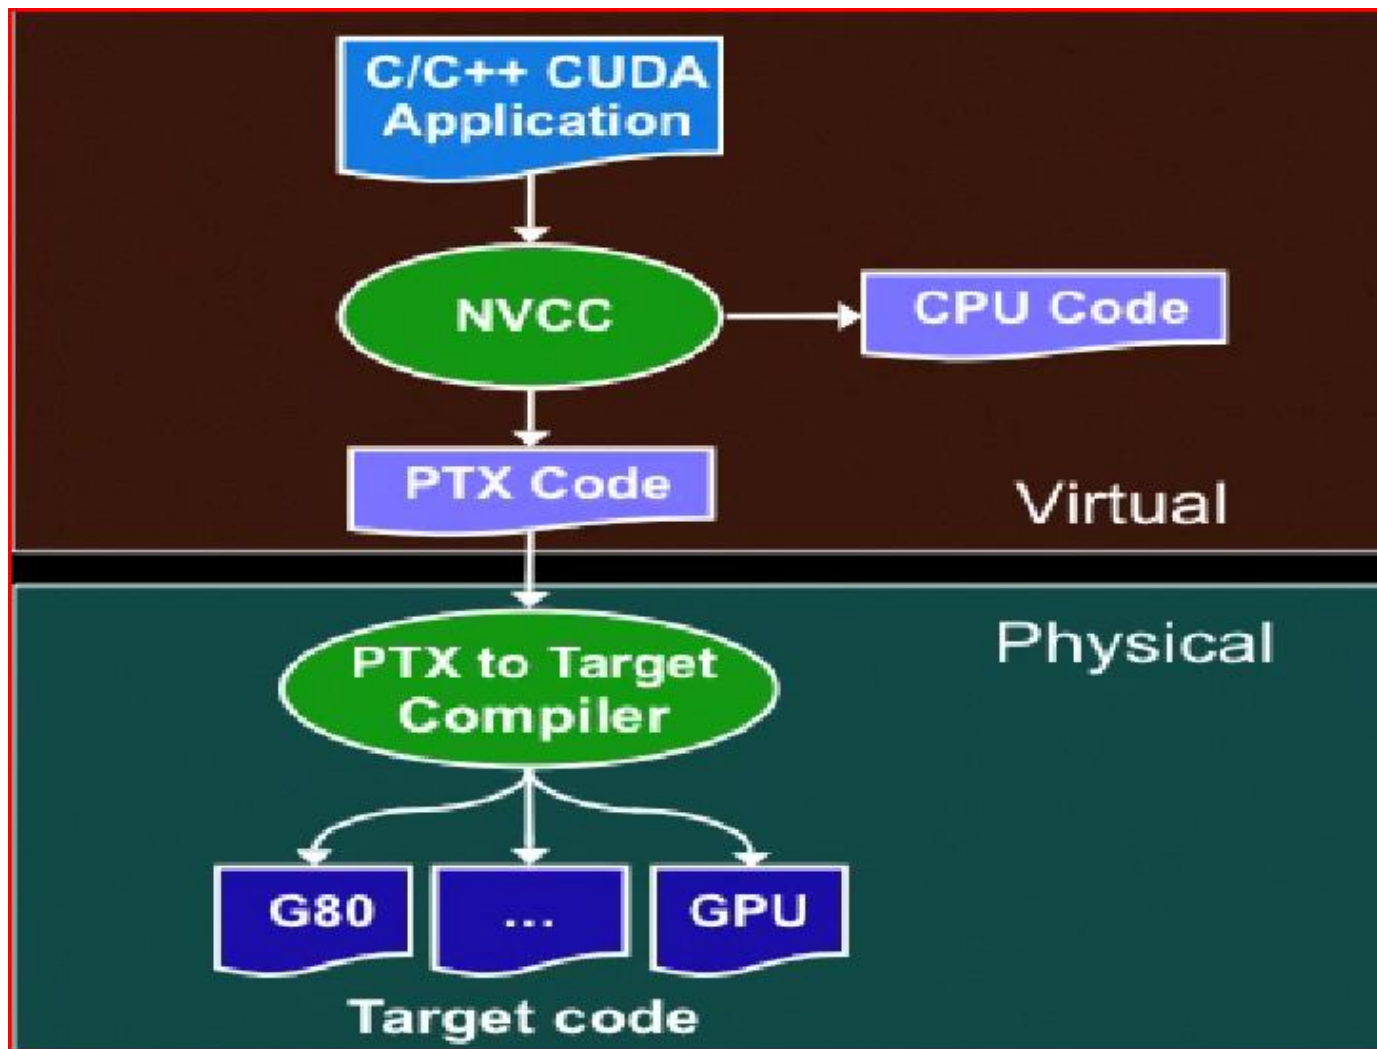
- **`dim3 gridDim;`**
 - Dimensions of the grid in blocks (at most 2D)
- **`dim3 blockDim;`**
 - Dimensions of the block in threads
- **`dim3 blockIdx;`**
 - Block index within the grid
- **`dim3 threadIdx;`**
 - Thread index within the block



*blockDim.x = 4
blockDim.y = 3
blockDim.z = 1*

CUDA Compile

*gridDim.x = 4
gridDim.y = 3
gridDim.z = 1*



CUDA Compile

nvcc <filename>.cu [-o <executable>]

- Builds release mode

nvcc -g <filename>.cu

- Builds debug mode
- Can debug host code but not device code

nvcc -deviceemu <filename>.cu

- Builds device emulation mode
- All code runs on CPU, no debug symbols

nvcc -deviceemu -g <filename>.cu

- Builds debug device emulation mode
- All code runs on CPU, with debug symbols

Hello World!

```
int main(void) {  
    printf("Hello World!\n");  
    return 0;  
}
```

Output:

- Standard C that runs on the host
 - NVIDIA compiler (nvcc) can be used to compile programs with no *device* code
- ```
$ nvcc
hello_world.cu
$ a.out
Hello World!
$
```
- Output* ←

# Hello World! with Device Code

```
__global__ void mykernel(void) {
}
```

```
int main(void) {
 mykernel<<<1,1>>>();
 printf("Hello World!\n");
 return 0;
}
```

- Two new syntactic elements...

# Hello World! with Device COde

```
mykernel<<<1,1>>>();
```

- Triple angle brackets mark a call from *host* code to *device* code
  - Also called a “kernel launch”
  - We’ll return to the parameters (1,1) in a moment
- That’s all that is required to execute a function on the GPU!

# Hello World! with Device Code

```
__global__ void mykernel(void) {
}
```

```
int main(void) {
 mykernel<<<1,1>>>();
 printf("Hello World!\n");
 return 0;
}
```


•mykernel() does nothing

cpu  
print it

Output:

```
$ nvcc
hello.cu
$ a.out
Hello World!
$
```

# Hello World! with Device Code



```
__global__ void mykernel(void) {
 printf("Hello World!\n");
}
```

```
int main(void) {
 mykernel<<<1,1>>>();
 return 0;
}
```

Output:

```
$ nvcc
hello.cu
$ a.out
Hello World!
$
```



Gpu  
print it

# Hello World! with Device Code

```
__global__ void mykernel(void) {
 printf("Hello World!\n");
}
```

```
int main(void) {
 mykernel<<<2,2>>>();
 return 0;
}
```

2 block  
2 thread

GPU

Output:

```
$ nvcc
hello.cu
```

```
$ a.out
```

```
Hello World!
Hello World!
Hello World!
Hello World!
$
```