

## Cuda Programming: Hello world:

- Cuda programming:
  - Kernels are C functions that can only be accessed from the GPU and must have void returns.
- Cuda built-in device variables:
  - gridDim, blockDim, blockIdx, threadIdx.

# Hello World! with Device Code

```
__global__ void mykernel(void) {  
    printf("Hello World!\n");  
}
```

```
int main(void) {  
    mykernel<<<1,1>>>();  
    return 0;  
}
```


## Output:

```
$ nvcc  
hello.cu  
$ a.out  
Hello World!  
$
```

## Cuda programming: sum of 2 arrays:

- GPU computing is about massive parallelism
- Memory management:
  - Device pointers point to GPU memory
  - Host pointers point to CPU memory
- Cuda API for handling device memory:
  - cudaMalloc(), cudaFree(), cudaMemcpy()
- Cuda threads:
  - A block can be split into parallel threads.
  - On the device, each thread can execute in parallel.

- Combining blocks and threads (Indexing):




- With M threads/block a unique index for each thread is given by:

```
int index = threadIdx.x + blockIdx.x * M;
```

## Formatting the grid as a Matrix


- dim3 grid (3,2);
- kernel<<<grid, 1>>>(...);



- ```
int index = blockIdx.x + blockIdx.y * gridDim.x;
```

# Formatting the grid as a Matrix

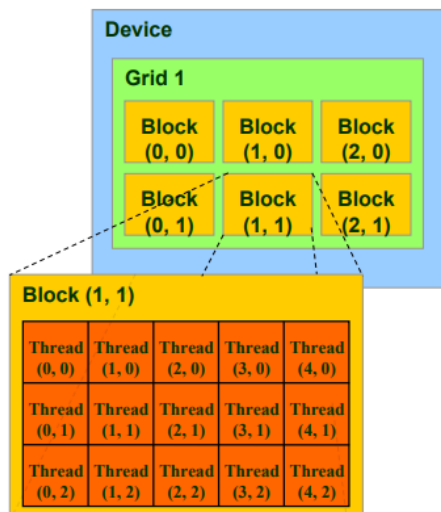
- `dim3 threads(5,3);`
- `kernel<<1, threads>>(...);`



- `int index = threadIdx.x + threadIdx.y * blockDim.x;`

# Formatting the grid as a Matrix

- `dim3 grid(3,2);`
- `dim3 block(5,3);`
- `kernel<<<grid, block>>>(...);`



- `int index = ( blockIdx.x + blockIdx.y * gridDim.x )  
* (blockDim.x * blockDim.y)  
+ threadIdx.y * blockDim.x + threadIdx.x;`

Cuda programming: Addition of 2 matrices:

- When adding matrices together, the `rowIndex` is usually `threadIdx.y` while the column index is usually `threadIdx.x`.
- The `y` is for rows, the `x` is for columns.
- So the addition code could look like this:

```
rowIndex = threadIdx.y;
colIndex = threadIdx.x;
```

```
c[rowIndex][colIndex] = a[rowIndex][colIndex] + b[rowIndex][colIndex];
```

- After indexing blocks and threads to make it parallel, the code looks like this:

## • Parallelized `add()` kernel

```
__global__ void add(int *a, int *b, int *c) {
    int rowIndex, colIndex;

    rowIndex = blockIdx.y * blockDim.y + threadIdx.y;
    colIndex = blockIdx.x * blockDim.x + threadIdx.x;

    c[rowIndex][colIndex] = a[rowIndex][colIndex] +
                             b[rowIndex][colIndex];
}
```

Running the kernel with `(N * N)` grid Blocks with `(N * N)` threads each.

# Matrix Addition on the Device: `add()`

- Parallelized `add()` kernel

```
__global__ void add(int *a, int *b, int *c, int width) {
    int rowIndex, colIndex, width;

    rowIndex = (blockIdx.y * blockDim.y + threadIdx.y) * width;
    colIndex = (blockIdx.x * blockDim.x + threadIdx.x) * width;

    for (int i=0; i < width; i++)
        for (int j=0; j < width; j++)
            c[rowIndex + i][colIndex + j] =
                a[rowIndex + i][colIndex + j] +
                b[rowIndex + i][colIndex + j];
}
```

- the same as the previous code but with width.

Cuda programming: product of matrices:

## Matrix multiplication

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 5 & 6 \\ 0 & 7 \end{bmatrix} = \begin{bmatrix} 1*5 + 2*0 & 1*6 + 2*7 \\ 3*5 + 4*0 & 3*6 + 4*7 \end{bmatrix} = \begin{bmatrix} 5 & 20 \\ 15 & 46 \end{bmatrix}$$

$$C[i][j] = \sum_{k=0}^n A[i][k] * B[k][j]$$

## Matrix Product on the Device: `add()`

- Parallelized `product()` kernel

```
__global__ void product(int *a, int *b, int *c, int n) {  
    int rowIndex, colIndex, width;  
  
    rowIndex = blockIdx.y * blockDim.y + threadIdx.y;  
    colIndex = blockIdx.x * blockDim.x + threadIdx.x;  
  
    for (int k=0; k < n; k++)  
        c[rowIndex][colIndex] += a[rowIndex][k] * b[k][colIndex];  
}
```

# Matrix Product on the Device: `add()`

- Parallelized `product()` kernel

```
__global__ void product(int *a, int *b, int *c, int width, int n) {
    int rowIndex, colIndex, width;

    rowIndex = (blockIdx.y * blockDim.y + threadIdx.y) * width;
    colIndex = (blockIdx.x * blockDim.x + threadIdx.x) * width;

    for (int i=rowIndex; i < rowIndex + width; i++)
        for (int j=colIndex; j < colIndex + width; j++)
            for (int k=0; k < n; k++)
                c[i][j] += a[i][k] * b[k][j];
}
```

- Same as the previous code but with width

Cuda programming: Querying the device:

cudaGetDeviceCount(int count) → returns the number of devices in the system.

char name[256] → an ASCII string identifying the device

size\_t totalGlobalMem → the amount of global memory on the device in bytes

size\_t sharedMemPerBlock → the maximum amount of shared memory a single block may use in bytes

int regsPerBlock → the number of 32-bit registers available per block

int warpSize → the number of threads in a warp

size\_t totalConstMem → the amount of available constant memory

int maxThreadsPerBlock → the maximum number of threads that a block may contain

int maxThreadsDim[3] → the maximum number of threads allowed along each dimension of a block

int maxGridSize[3] → the number of blocks allowed along each dimension of a grid

int multiProcessorCount → the number of multiprocessors on the device

int concurrentKernels → a Boolean value representing whether the device supports executing multiple kernels within the same context simultaneously

```
int main(void) {
    int count;
    cudaDeviceProp prop;

    cudaGetDeviceCount( &count);
    for (int i=0; i< count; i++) {
        cudaGetDeviceProperties( &prop, i );
        //Do something with our device's properties
    }
    return 0;
}
```



```

int main(void) {
    cudaDeviceProp prop;
    int dev;

    cudaGetDevice( &dev );
    printf( "ID of current CUDA device: %d\n", dev );
    memset( &prop, 0, sizeof( cudaDeviceProp ) );
    prop.major = 1;
    prop.minor = 3;
    cudaChooseDevice( &dev, &prop );
    printf( "ID of device closest to revision 1.3: %d\n", dev );
    cudaSetDevice( dev );
    return 0;
}

```

© NVIDIA 2013

Parallel sorting algorithms:

Odd-even sort:

- The odd-even sort is based on the bubble-sort technique
- Adjacent pairs of items in an array are exchanged if they are found to be out of order
- Total running time is  $O(\log^2 N)$
- Code for this technique:

## Odd-Even Sort Algorithm

---

### Algorithm 1 Odd Even Sort

---

```

for  $k = 1 \rightarrow N/2$  do
  do parallel
    if  $i > i + 1 \ \forall \ i \% 2 \neq 0$  then
      swap  $i, i + 1$ 
    end if
  end parallel
  do parallel
    if  $i > i + 1 \ \forall \ i \% 2 == 0$  then
      swap  $i, i + 1$ 
    end if
  end parallel
end for

```

---

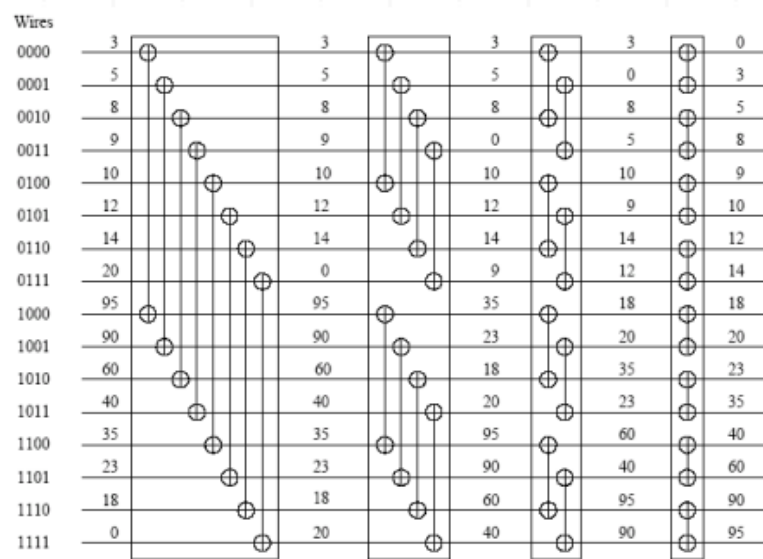


### Bitonic sort:

- A bitonic sorting network sorts  $n$  elements in  $O(\log^2 n)$  time
- Bitonic sequence has two tones  $\rightarrow$  increasing and decreasing.


- The kernel of the network is the rearrangement of a bitonic sequence into a sorted

## Sorting Networks: Bitonic Sort



A bitonic merging network for  $n = 16$ . The entire figure represents a  $\oplus\text{BM}[16]$  bitonic merging network. The network takes a bitonic sequence and outputs it in sorted order.

sequence



A schematic representation of a network that converts an input sequence into a bitonic sequence. In this example,  $\oplus\text{BM}[k]$  and  $\ominus\text{BM}[k]$  denote bitonic merging networks of input size  $k$  that use  $\oplus$  and  $\ominus$  comparators, respectively. The last merging network ( $\oplus\text{BM}[16]$ ) sorts the input. In this example,  $n = 16$ .

|                              |                   |                |                |                                               |                                   |   |   |   |   |
|------------------------------|-------------------|----------------|----------------|-----------------------------------------------|-----------------------------------|---|---|---|---|
|                              | 1                 |                |                | 2                                             |                                   |   | 3 |   |   |
| Step                         | 1                 | 2              | 3              | 4                                             | 5                                 | 6 | 7 | 8 | 9 |
| phase                        | 1                 | 1              | 2              | 1                                             | 2                                 | 1 | 2 | 3 | 4 |
| $n = \text{sequence length}$ | $2^1 = 2$         | 4              | 2              | $n = 2^i / 2^{j-1}$                           | $n = 2^{i-j+1}$                   |   | 8 | 4 | 2 |
| offset (shift)               | $\frac{n}{2} = 1$ | 2              | 1              | $\frac{n}{2} = \frac{2^{i-j+1}}{2} = 2^{i-j}$ |                                   |   |   |   |   |
| involved threads             | even              | $Tid \% 4 < 2$ | $Tid \% 2 < 1$ | $Tid \% n < \frac{n}{2}$                      | $Tid \% 2^{i-j+1} < 2^{i-j}$      |   |   |   |   |
| operation                    |                   |                |                | $\oplus$ if $(Tid / 2^i)$ is even             | $\ominus$ if $(Tid / 2^i)$ is odd |   |   |   |   |

Dynamic parallelism:

C  $\rightarrow$  Compute


U  $\rightarrow$  Unified

D → Device

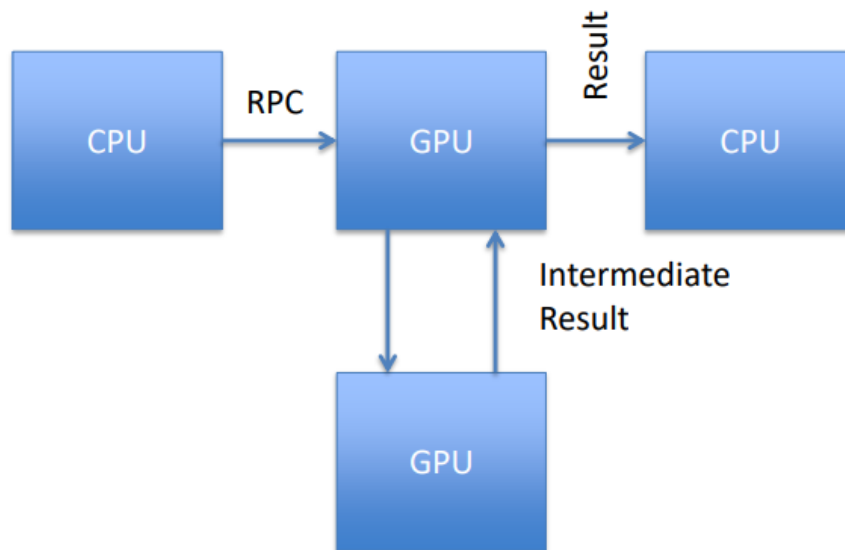
A → Architecture

- Compute unified architecture:
  - Hybrid CPU/GPU code
  - Low latency code is running on the CPU while the high latency code is running on the GPU
- Types of parallelism:
  - Task parallelism:
    - Problem is divided to tasks, which are processed independently
  - Data parallelism:
    - Same operation is performed over many data items
  - Pipeline parallelism:
    - Data are flowing through a sequence of stages, which operate concurrently
- Parallelism in GPU:
  - Data parallelism:
    - The same kernel is executed by many threads
    - Thread process one data item
  - Limited task parallelism:
    - Multiple kernels executed simultaneously
    - At most as many kernels as SMP
  - We don't have:
    - Any means of kernel-wide synchronization
    - Any guarantees that two blocks/kernels will actually run concurrently
- What is dynamic parallelism:
  - The ability to launch new kernels from the GPU:
    - Dynamically
      - Based on run time
    - Simultaneously:
      - From multiple threads at once
    - Independently:
      - Each thread can launch a different grid

# Execution Model (Overview)



*Fermi: Only CPU can generate GPU work*



*Kepler: GPU can generate work for itself*

- Dynamic parallelism:
  - Allows program flow to be controlled by the GPU
  - Allows recursion and subdivision of problems
  - Interesting data is not uniformly distributed
  - Dynamic parallelism can launch additional threads in interesting areas
  - Allows higher resolution in critical areas without slowing down others
- Problem cases:
  - Unsuitable problems for GPUs:
    - Processing irregular data structures
    - Regular structures with irregular processing workload
    - Iterative tasks with explicit synchronization
    - Pipeline-oriented tasks with many simple stages
  - Solutions:
    - Iterative kernel execution
    - Mapping irregular structures to regular grids
    - 2-phase algorithms
- Dynamic parallelism purpose:
  - The device does not need to synchronize with host to issue new work to the device
  - Irregular parallelism may be expressed more easily

- Without dynamic parallelism:
  - Data travels back and forth between the CPU and GPU many times
  - This is because of the inability of the GPU to create more work on itself depending on the data
- With dynamic parallelism:
  - GPU can generate work on itself based on intermediate results, without the involvement of CPU
  - Permits dynamic run time decisions
  - Leaves the CPU free to do other work
- How does dynamic parallelism work?
  - Portions of CUDA runtime are ported to device side:
    - Kernel execution
    - Device synchronization
    - Streams, events, and async memory operations
  - Kernel launches are asynchronous:
    - No guarantee the child kernel starts immediately
    - Synchronization points may cause context switch
  - Block-wise locality of resources:
    - Streams and events are shared in thread block

## Dynamic Parallelism

### • Example

```
__global__ void child_launch(int *data) {
    data[threadIdx.x] = data[threadIdx.x]+1;
}

__global__ void parent_launch(int *data) {
    data[threadIdx.x] = threadIdx.x;
    __syncthreads();
    if (threadIdx.x == 0) {
        child_launch<<< 1, 256 >>>(data);
        cudaDeviceSynchronize();
    }
    __syncthreads();
}

void host_launch(int *data) {
    parent_launch<<< 1, 256 >>>(data);
}
```

Thread 0 invokes a grid of child threads

Synchronization does not have to be invoked by all threads

Device synchronization does not synchronize threads in the block


- More on dynamic parallelism:
  - Scheduling can be controlled by streams
  - Launched kernels may execute out-of-order within a stream
  - Named streams can guarantee concurrency
  - Nested dependencies:

- cudaDeviceSynchronize()
- can be used inside a kernel
- synchronizes all launches by any kernel in block
- does not imply \_\_syncthreads()
- kernel launch implied memory sync operation
- child sees state at time of launch
- parent sees child writes after sync
- local and shared memory are private, cannot be shared with children

## Optimizing parallel reduction in CUDA:

- Parallel reduction:
  - Common and important data parallel primitive
  - Easy to implement in CUDA
  - Serves as a great optimization example
  - It has a tree-based approach used within each thread block

### ● Tree-based approach used within each thread block



- Problem: global synchronization:
  - CUDA has no global synchronization because it's expensive to build in hardware for GPUs with high processor count
  - The solution is to decompose into multiple kernels:
    - Kernel launches serves as a global synchronization point
    - Kernel launch has negligible HW overhead, low SW overhead



- Reduction 1 → interleaved addressing with divergent branching:

## Reduction #1: Interleaved Addressing



```
__global__ void reduce0(int *g_idata, int *g_odata) {
    extern __shared__ int sdata[];


    // each thread loads one element from global to shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();

    // do reduction in shared mem
    for(unsigned int s=1; s < blockDim.x; s *= 2) {
        if (tid % (2*s) == 0) {
            sdata[tid] += sdata[tid + s];
        }
        __syncthreads();
    }

    // write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

7

## Parallel Reduction: Interleaved Addressing



8

- The problem with this method is that highly divergent warps are inefficient, and the % operator is very slow

- Reduction 2 → interleaved addressing with bank conflicts:

## Reduction #2: Interleaved Addressing



Just replace divergent branch in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {
    if (tid % (2*s) == 0) {
        sdata[tid] += sdata[tid + s];
    }
    __syncthreads();
}
```


With strided index and non-divergent branch:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {
    int index = 2 * s * tid;

    if (index < blockDim.x) {
        sdata[index] += sdata[index + s];
    }
    __syncthreads();
}
```

11

## Parallel Reduction: Interleaved Addressing



**New Problem: Shared Memory Bank Conflicts**

12

- Reduction 3 → sequential addressing:

## Reduction #3: Sequential Addressing



Just replace strided indexing in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {
    int index = 2 * s * tid;

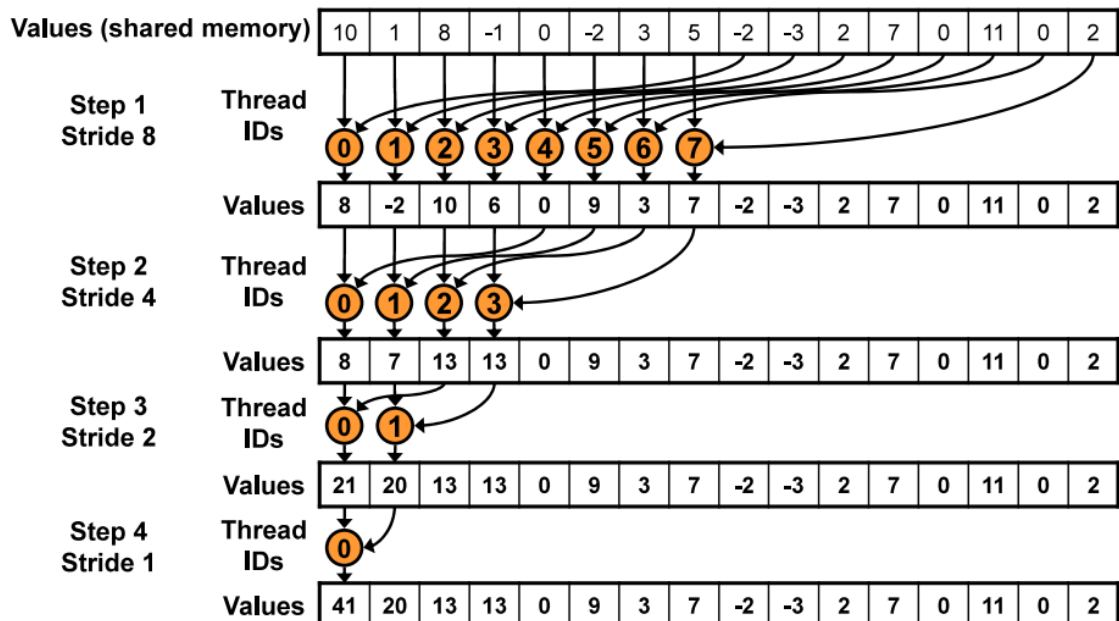
    if (index < blockDim.x) {
        sdata[index] += sdata[index + s];
    }
    __syncthreads();
}
```

With reversed loop and threadID-based indexing:

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
    if (tid < s) {
        sdata[tid] += sdata[tid + s];
    }
    __syncthreads();
}
```

15

## Parallel Reduction: Sequential Addressing



Sequential addressing is conflict free

14

- Problem with reduction 3:
  - o Idle threads
    - Half of the threads are idle on the first loop iteration
- Reduction 4 → first add during load:

## Reduction #4: First Add During Load



Halve the number of blocks, and replace single load:

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
sdata[tid] = g_idata[i];
__syncthreads();
```

With two loads and first add of the reduction:

```
// perform first level of reduction,
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

- We could have instruction bottleneck
- So, the next improvement (reduction 5) will be unrolled loops
- Unrolling the last warp:
  - o As reduction proceeds, the number of active threads decrease
  - o Instructions are SIMD synchronous within a warp

- Reduction 5 → unroll the last warp:

## Reduction #5: Unroll the Last Warp



```
__device__ void warpReduce(volatile int* sdata, int tid) {
    sdata[tid] += sdata[tid + 32];
    sdata[tid] += sdata[tid + 16];
    sdata[tid] += sdata[tid + 8];
    sdata[tid] += sdata[tid + 4];
    sdata[tid] += sdata[tid + 2];
    sdata[tid] += sdata[tid + 1];
}
```

**IMPORTANT:**  
For this to be correct,  
we must use the  
"volatile" keyword!

```
// later...
for (unsigned int s=blockDim.x/2; s>32; s>>=1) {
    if (tid < s)
        sdata[tid] += sdata[tid + s];
    __syncthreads();
}

if (tid < 32) warpReduce(sdata, tid);
```

**Note:** This saves useless work in *all* warps, not just the last one!

Without unrolling, all warps execute every iteration of the for loop and if statement

22

- Performance for all 5 reductions:

## Performance for 4M element reduction



|                                                                        | Time (2 <sup>22</sup> ints) | Bandwidth   | Step Speedup | Cumulative Speedup |
|------------------------------------------------------------------------|-----------------------------|-------------|--------------|--------------------|
| <b>Kernel 1:</b><br>interleaved addressing<br>with divergent branching | 8.054 ms                    | 2.083 GB/s  |              |                    |
| <b>Kernel 2:</b><br>interleaved addressing<br>with bank conflicts      | 3.456 ms                    | 4.854 GB/s  | 2.33x        | 2.33x              |
| <b>Kernel 3:</b><br>sequential addressing                              | 1.722 ms                    | 9.741 GB/s  | 2.01x        | 4.68x              |
| <b>Kernel 4:</b><br>first add during global load                       | 0.965 ms                    | 17.377 GB/s | 1.78x        | 8.34x              |
| <b>Kernel 5:</b><br>unroll last warp                                   | 0.536 ms                    | 31.289 GB/s | 1.8x         | 15.01x             |

## Reasoning about performance:

- Parallelism vs Performance:
  - Ideally a problem that takes  $T$  time on a single processor can execute in  $T/P$  time on  $P$  processors
  - The challenge of meeting  $T/P$  goal becomes more difficult as  $P$  increases
  - There are certain cases where the  $P$  processors produce shorter execution time than  $T/P$
- First code, serial 3s count program:

### Serial 3s Count Program

```
public class Count3S {  
    int array[];  
    int count;  
  
    public int count3s() {  
        count = 0;  
        for (int i = 0; i < array.length; i++) {  
            if (array[i] == 3)  
                count ++;  
        }  
        return count;  
    }  
}
```

- Second code, parallel 3s count program (first one):

## Parallel 3s Count Program

```
public class Count3sParallel1 implements Runnable {
    int array[];
    int count, nbThread;
    Thread t;
    LinkedList<Integer> threadIds = new LinkedList<Integer>();

    public void count3s() {
        count = 0;
        for (int i=0; i < nbThread; i++) {
            t = new Thread(this);
            threadIds.add(new Integer(i));
            t.start();
        }
    }

    public void run() {
        int depth = (array.length / nbThread);
        int start = threadIds.poll().intValue() * depth;
        int end = start + depth;

        for (int i = start; i < end; i++) {
            if (array[i] == 3)
                count ++;
        }
    }
}
```

- The problem with this program is that the threads can all access the array and increment the count variable at the same time, which will result in inaccurate results.

- Third code, parallel 3s count program (second one):

```
public class Count3SParallel2 implements Runnable {
    int array[];
    int count, nbThread;
    Thread t;
    LinkedList<Integer> threadIds = new LinkedList<Integer>();
    Mutex m = new Mutex();

    public void count3s() {
        count = 0;
        for (int i=0; i < nbThread; i++) {
            t = new Thread(this);
            threadIds.add(new Integer(i));
            t.start();
        }
    }

    public void run() {
        int depth = (array.length / nbThread);
        int start = threadIds.poll().intValue() * depth;
        int end = start + depth;

        for (int i = start; i < end; i++) {
            if (array[i] == 3) {
                try {
                    m.acquire();
                    count++;
                    m.release();
                } catch (InterruptedException e) {
                    e.printStackTrace();
                }
            }
        }
    }
}
```

- Here, we fixed the issue of threads accessing at the same time by mutual exclusion (Mutex m) which will act as a key for threads to not access at the same time
- However, this will bring a lot of idle time on the table so further modifications are required



- Fourth code, parallel 3s count program (third one):

```
public class Count3sParallel3 implements Runnable {
    int array[];
    int count, nbThread;
    Thread t;
    LinkedList<Integer> threadIds = new LinkedList<Integer>();
    Mutex m = new Mutex();

    public void count3s() {
        count = 0;
        for (int i=0; i < nbThread; i++) {
            t = new Thread(this);
            threadIds.add(new Integer(i));
            t.start();
        }
    }

    public void run() {
        int depth = (array.length / nbThread);
        int start = threadIds.poll().intValue() * depth;
        int end = start + depth;
        int localCount = 0;

        for (int i = start; i < end; i++) {
            if (array[i] == 3)
                localCount++;
        }
        try {
            m.acquire();
            count += localCount;
            m.release();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}
```

- We fixed the issue of mutex by adding a local variable that will speed up the addition and threads don't have to keep going in and out of the try block to add
- Another problem is posed here, which is copying data from local memory to global memory which is slow

- Fifth code, parallel 3s count (fourth one):

```
public class Count3sParallel4 implements Runnable {
    int array[];
    int count, nbThread;
    Thread t;
    LinkedList<Integer> threadIds = new LinkedList<Integer>();
    Mutex m = new Mutex();
    int localCounts[ ];

    public void count3s() {
        count = 0;
        for (int i=0; i < nbThread; i++) {
            t = new Thread(this);
            threadIds.add(new Integer(i));
            t.start();
        }
    }

    public void run() {
        int id = threadIds.poll().intValue();
        int depth = (array.length / nbThread);
        int start = id * depth;
        int end = start + depth;

        for (int i = start; i < end; i++) {
            if (array[i] == 3)
                localCounts[id]++;
        }
        try {
            m.acquire();
            count += localCounts[id];
            m.release();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}
```

- Here, we fixed the issue of the previous run by making the local count a global array of 0 and each thread will increment the value of its slots in the array and then in the try block they only copy once.
- Threads:
  - A thread consists mainly of:
    - A program code
    - A call stack
    - And some modest amount of thread specific data
  - Threads share access to memory:
    - Communicate by reading from or writing to memory that is visible to all
    - Shared-memory parallel programming
- Processes:
  - A process is a thread that also has its own private address space
  - Processes communicate by passing messages
  - Processes have much more associated state than threads
- Latency and throughput:
  - Latency refers to the amount of time it takes to complete a given unit of work
  - Throughput refers to the amount of work that can be completed per unit of time
  - Parallelism can be exploited to improve throughput
  - Parallelism can be exploited to hide/reduce latency

- Source of performance loss:
  - Overhead
  - Non-parallelizable computation
  - Idle processors
  - Contention of resources
- Overhead:
  - Any cost that is incurred in the parallel solution but not in the serial solution
  - Sources of overhead:
    - Communication
    - Synchronization
    - Computation
    - Memory
- Non-parallelizable code:
  - Amdahl's Law:
    - If  $1/S$  of a computation is inherently sequential
    - The maximum performance improvement is a factor of  $S$
    - $T_p = (1/S * T_s) + ((1-1/S) * T_s/P)$
- Idle time:
  - Idle time is a consequence of synchronization and communication
  - Ideally, all processors are working all the time
- Contention:
  - Contention is the degradation of system performance caused by computation
  - Contention can lead to slow down performance
  - Lock contention can reduce performance