Remaining Time: 25 minutes, 26 seconds.

Question Completion Status:

QUESTION 1

Successive calls of different kernels are:

- ○ Performed concurrently.
- ○ Permed simultaneously.
- ◉ Performed sequentially.
- ○ Performed in parallel

QUESTION 2

Which of following statements allows to get the nu

- ○ int n;

  n = cudaGetDevice ();

Click Save and Submit to save and submit. Click Save

int n;
cudaGetDeviceCount(&n);

QUESTION 3

Access to Local Memory is slow because

- It has a relatiively large size.
- It is mounted on ship.
- It is mounted on the DRAM.
- it is used in read and write modes.

QUESTION 4

Click Save and Submit to save and submit.

QUESTION 2

Which of following statements allows to get the number of available devices:

- int n;
  n = cudaGetDevice ();
- int n;
  n = cudaGetDeviceCount();
- int n;
  cudaGetDevice(&n);
- int n;
  cudaGetDeviceCount(&n);

QUESTION 5

GPGPU is the abbreviation of:

- (●) General Purpose GPU
- ( ) Global Processing GPU
- ( ) General Processing GPU
- ( ) Global Purpose GPU

QUESTION 4

In CUDA, threads _____.

- [x] are organized in 1-D, 2-D or 3-D blocks.
- [x] of the same block run in groups of 32 called warps.
- [ ] can be synchronized with threads of other blocks.
- [ ] can share data with each others using the shared memory.

QUESTION 7

A kernel is performed by a grid of thread-blocks. A Grid could be :

- [x] 2-D array of thread-blocks.
- [ ] 1-D, 2-D or 3-D array of thread-blocks.
- [ ] 3-D array of thread-blocks.
- [x] 1-D array of thread-blocks.

QUESTION 6

1 poir

Which of the following statements allows to identify the currently active device that wil run the kernel code.

- cudaChooseDevice(...);
- cudaGetDevice(...);
- cudaGetDeviceProperties(...)
- cudaSetDevice(...);

QUESTION 9

Dynamic parallelism allows :

- ○ the device to run successive calls of kernels simulateously.
- ○ the device to run repetitive (iterative) calls of the same kernel simulateously.
- ○ the host to lauch several kernels simulatenously.
- ◉ the device to launch new kernels at run time.

Question Completion Status:

1-D, 2-D or 3-D array of thread-blocks.

3-D array of thread-blocks.

1-D array of thread-blocks.

QUESTION 8

Registers have fast access because:

they are on the DRAM.

they are cached.

they have small size.

they are mounted on chip.

QUESTION 9

Click Save and Submit to save and submit. Click

QUESTION 11

Dynamic parallelism is suitable :

- [x] when we are processing irregular data structures such as trees and graphs.
- [x] when the work-load is irregular while processing a regular data structure.
- [ ] when the host is launching repetitively the same kernel.
- [ ] when the host is launching several kernels simultaneously.

QUESTION 10

Using CUDA:

- [ ] Low latency code is running on the GPU device.
- [x] High latency code is running on the GPU device.
- [x] Low latency code is running on the CPU.
- [ ] High latency code is running on the host.