

PARALLEL PROCESSING

Mohammed Alabdulkareem

kareem@ksu.edu.sa

Office 2247


PROGRAMMING DISTRIBUTED MEMORY MACHINE



- Shared memory machines have faster communication time than distributed memory machine but it is less scalable in terms of number of processors.
- Distributed memory machines are more scalable in number of processors. But with more communication time.
- Hence the speed of the communication network need to be fast enough to close this gap.


PROGRAMMING DISTRIBUTED MEMORY MACHINE

- Distributed memory machines can not communicate through shared memory.
- Exchanging data and communications are done using communication network.

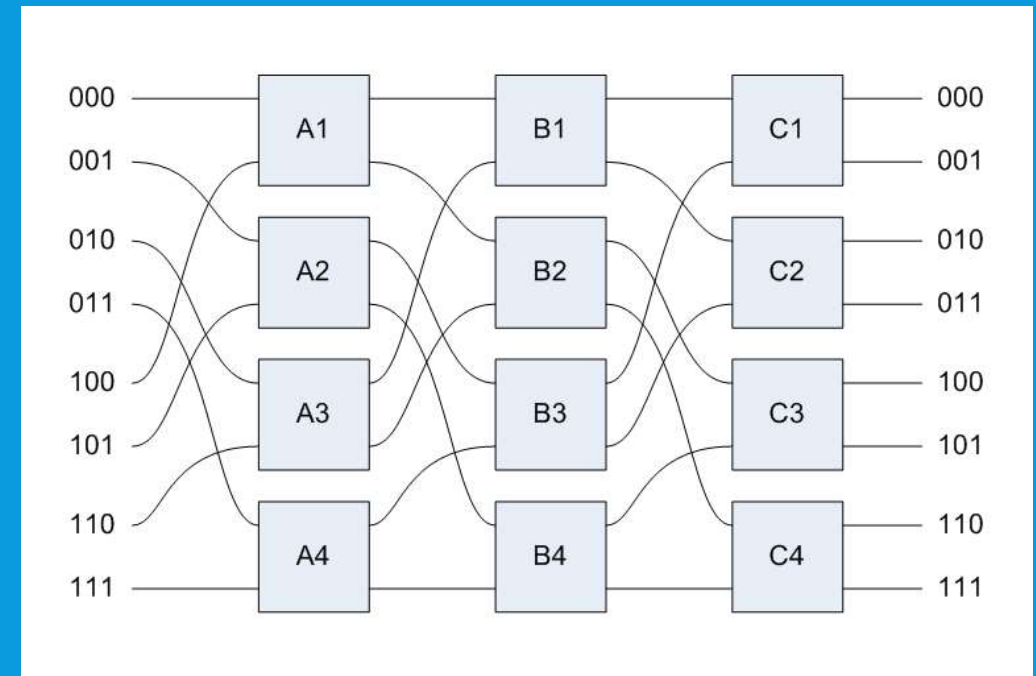


PROGRAMMING DISTRIBUTED MEMORY MACHINE

Interconnection networks (examples)



Hypercube



Omega

PROGRAMMING DISTRIBUTED MEMORY MACHINE

Supercomputers with interconnection network



PROGRAMMING DISTRIBUTED MEMORY MACHINE



- In distributed memory machines cooperation among processes implies the data exchange using message passing (communication overhead).
- The overall execution time is consequently a sum of computation and communication time.
- Algorithms with only local communication between neighboring processors are faster and more scalable than the algorithms with the global communication among all processors.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI is **M**essage **P**assing Interface.
- MPI is introduced from the practical point of view, with a set of basic operations that enable implementation of parallel programs.
- The MPI library interface is a specification, not an implementation.
- The MPI is not a language, and all MPI operations are expressed as functions, subroutines, or methods, according to the appropriate language bindings for C and Fortran.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI standard includes:
 - Process creation and management.
 - Language bindings for C and Fortran.
 - Point-to-point and collective communications.
 - Group and communicator concepts.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI program consists of autonomous processes that are able to execute their own code in the sense of Multiple Instruction Multiple Data (MIMD) paradigm.
- The processes communicate via calls to MPI communication operations independently of operating system.
- Any MPI program should have operations to initialize execution environment and to control starting and terminating procedures of all generated processes.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI processes can be collected into **groups** of specific *size* that can communicate in its own environment where each message sent in a context must be received only in the same context.
- A process group and context together form an MPI communicator.
- A process is identified by its rank in the group associated with a communicator.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI has default communicator `MPI_COMM_WORLD`
- The group associated with this communicator includes all initial processes with default context.
- MPI operations return a status of the execution success; in C routines it is returned as the value of the function.
- The MPI standard *does not* specify how **st dout** from different nodes should be collected for printing at the originating process.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI_SEND (buf, count, datatype, dest, tag, comm)
- The basic MPI communication is characterized by two fundamental MPI operations MPI_SEND and MPI_RECV.
- MPI_SEND is blocking call it will not complete until there is matching MPI_RECV on the other receiving process.
- The MPI_RECV will empty the input send buffer *buf* of matching MPI_SEND.
- MPI_SEND will return when the message data has been delivered to the communication system and the send buffer *buf* of the sender process source can be reused.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- `MPI_SEND (buf, count, datatype, dest, tag, comm)`
 - *buf* is a pointer to the send buffer
 - *count* is the number of data items
 - *datatype* is the type of data items (more about data types in MPI)
 - *dest* the rank of receiver process (similar to thread number)
 - *tag* a message tag to distinguish between different messages for the same receiver process
 - *comm* the communicator.

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI_RECV (buf, count, datatype, source, tag, comm, status)
 - *buf* is a pointer to the output buffer
 - *count* is the number of data items to be received
 - *datatype* is the type of data items
 - *source* the rank of sending process
 - *tag* a message tag
 - *comm* the communicator.
 - *status* contains further information in case of error

PROGRAMMING DISTRIBUTED MEMORY MACHINE

MPI datatype	C equivalent
MPI_SHORT	short int
MPI_INT	int
MPI_LONG	long int
MPI_LONG_LONG	long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	char

- Wildcard may be used at the destination side

MPI_RECV (buf, count, datatype, source, tag, comm, status)

- *Source* MPI_ANY_SOURCE
- *Tag* MPI_ANY_TAG

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- **MPI_Init** initializes the MPI execution environment.
- **MPI_Comm_size(MPI_COMM_WORLD, &size)** returns size, which is the number of started processes.
- **MPI_Comm_rank(MPI_COMM_WORLD, &rank)** that returns rank or the ID of each process.
- **MPI_Finalize** exits the MPI

PROGRAMMING DISTRIBUTED MEMORY MACHINE

```
#include <stdio.h>
#include "mpi.h"
int main( argc, argv )
int argc;
char **argv;
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    printf( "Hello world from process %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

PROGRAMMING DISTRIBUTED MEMORY MACHINE



```
% mpicc -o helloworld helloworld.c  
% mpirun -np 4 helloworld  
Hello world from process 0 of 4  
Hello world from process 3 of 4  
Hello world from process 1 of 4  
Hello world from process 2 of 4  
%
```

PROGRAMMING DISTRIBUTED MEMORY MACHINE



```
int world_rank;

MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

int world_size;

MPI_Comm_size(MPI_COMM_WORLD, &world_size);

int number;

if (world_rank == 0) {

    number = -1;

    MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);

} else if (world_rank == 1) {

    MPI_Recv(&number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    printf("Process 1 received number %d from process 0\n", number);

}
```

PROGRAMMING DISTRIBUTED MEMORY MACHINE

```
MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
```

- *Buf* &number
- *count* 1
- *datatype* MPI_INT
- *dest* 1
- *tag* 0
- *comm* MPI_COMM_WORLD.

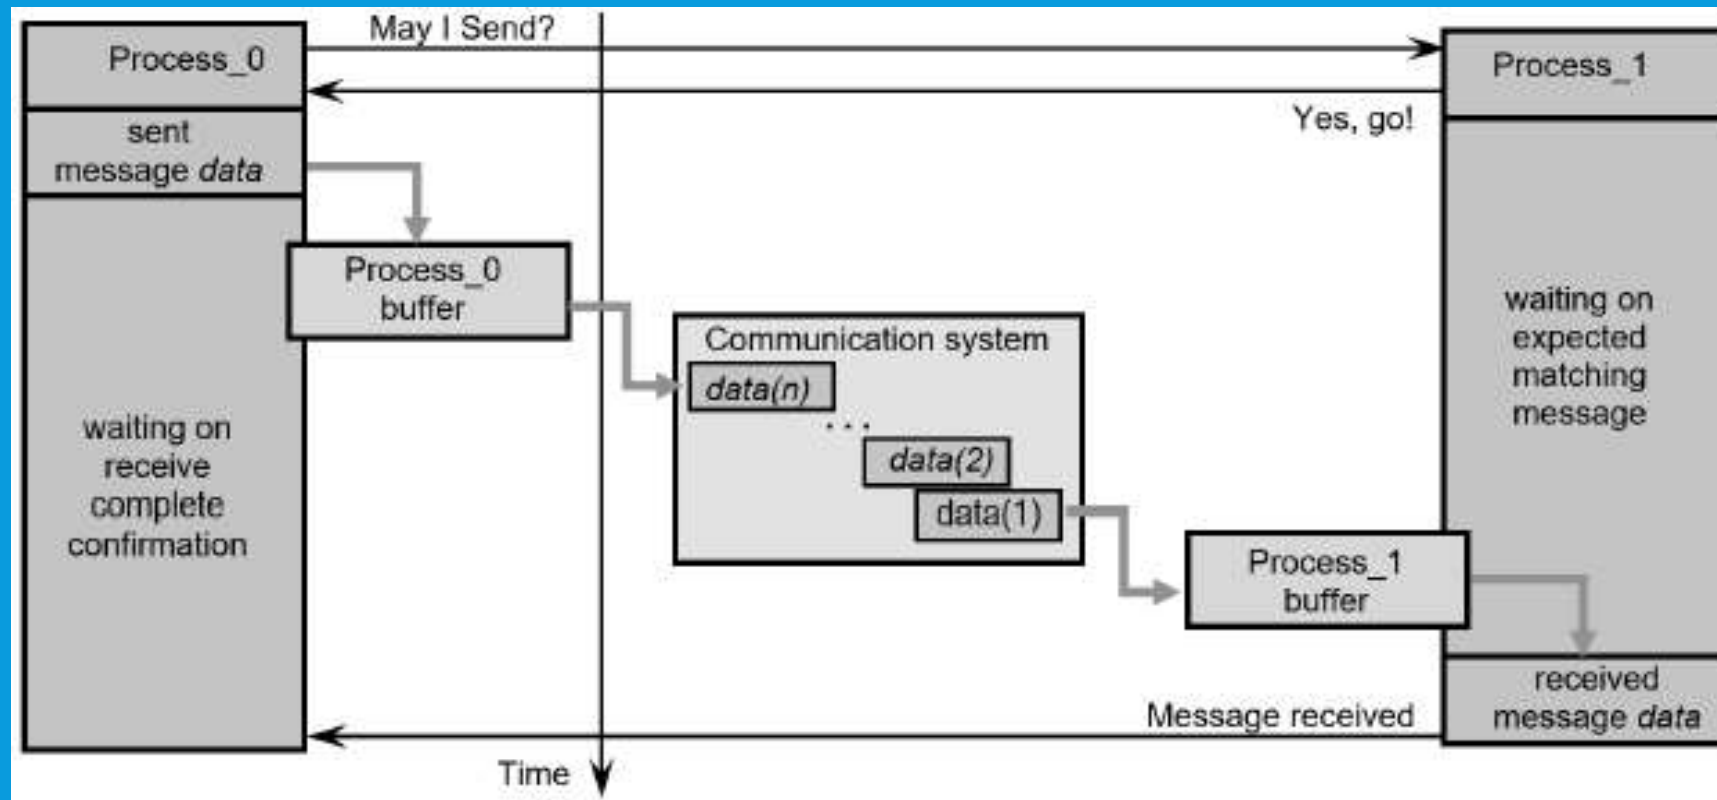
PROGRAMMING DISTRIBUTED MEMORY MACHINE



```
mpirun -n 2 ./send_recv  
Process 1 received number -1 from process 0
```

PROGRAMMING DISTRIBUTED MEMORY MACHINE

How it works



PROGRAMMING DISTRIBUTED MEMORY MACHINE



MPI_RECV (buf, count, datatype, source, tag, comm, status)

- The number of received data items of datatype must be equal or fewer as specified by count. ($\text{count} \geq 0$).
- Receiving more data items results in an error and the status value return more information.

PROGRAMMING DISTRIBUTED MEMORY MACHINE

- MPI_SENDRECV (sendbuf, sendcount, sendtype, dest, sendtag, recvbuf, recvcount, recvtype, source, recvtage, comm, status)



PROGRAMMING DISTRIBUTED MEMORY MACHINE



- MPI_SENDRECV (sendbuf, sendcount, sendtype, dest, sendtag, recvbuf, recvcount, recvtype, source, recvtag, comm, status)
- MPI_SENDRECV combines a sending of message to destination process **dest** and a receiving of another message from process **source**, in a single call in sender and receiver process.

Measuring time

```
double start, finish;  
start = MPI_Wtime ();  
... //MPI program segment to be measured  
finish = MPI_Wtime ();  
printf ("Elapsed time is %f\n", finish - start);
```

PROGRAMMING DISTRIBUTED MEMORY MACHINE



- The elapsed time (wall-clock) between two points in an MPI program can be measured by using operation `MPI_WTIME ()`.
- The time measured in seconds.

MPI_BARRIER (comm)

- This operation is used to synchronize the execution of a group of processes specified within the communicator **comm**.
- When a process reach this operation it waits until all other processes reach the **MPI_BARRIER**.
- The barrier is a simple way of separating two phases of a computation to ensure that messages generated in different phases do not interfere.

PROGRAMMING DISTRIBUTED MEMORY MACHINE

MPI_BARRIER (comm)

