



# FAST PARALLEL SORTING ALGORITHMS ON GPUS

# Introduction

*Odd- Even sort*

*Rank sort*

*Bitonic sort*



# Odd-Even Sort

1

# Odd-Even Sort

- The odd-even sort is a parallel sorting algorithm and is based on bubble-sort technique.
- Adjacent pairs of items in an array are exchanged if they are found to be out of order.
- Total running time for this technique is  $O(\log 2N)$ .

# Odd-Even Sort Algorithm

---


## Algorithm 1 Odd Even Sort


---

```
for  $k = 1 \rightarrow N/2$  do  
  do parallel  
    if  $i > i + 1 \vee i \% 2 \neq 0$  then  
      swap  $i, i + 1$   
    end if  
  end parallel  
  do parallel  
    if  $i > i + 1 \vee i \% 2 == 0$  then  
      swap  $i, i + 1$   
    end if  
  end parallel  
end for
```

---

# Odd-Even Sort Algorithm





Bitonic Sort

3

## Sorting Networks: Bitonic Sort

- A bitonic sorting network sorts  $n$  elements in  $\Theta(\log^2 n)$  time.
- A bitonic sequence has two tones - increasing and decreasing, or vice versa. Any cyclic rotation of such networks is also considered bitonic.
- $\langle 1, 2, 4, 7, 6, 0 \rangle$  is a bitonic sequence, because it first increases and then decreases.  $\langle 8, 9, 2, 1, 0, 4 \rangle$  is another bitonic sequence.
- The kernel of the network is the rearrangement of a bitonic sequence into a sorted sequence.



## Sorting Networks: Bitonic Sort

- Let  $s = \langle a_0, a_1, \dots, a_{n-1} \rangle$  be a bitonic sequence such that  $a_0 \leq a_1 \leq \dots \leq a_{n/2-1}$  and  $a_{n/2} \geq a_{n/2+1} \geq \dots \geq a_{n-1}$ .

- Consider the following subsequences of  $s$ :

$$s_1 = \langle \min\{a_0, a_{n/2}\}, \min\{a_1, a_{n/2+1}\}, \dots, \min\{a_{n/2-1}, a_{n-1}\} \rangle$$

$$s_2 = \langle \max\{a_0, a_{n/2}\}, \max\{a_1, a_{n/2+1}\}, \dots, \max\{a_{n/2-1}, a_{n-1}\} \rangle$$

(1)

Note that  $s_1$  and  $s_2$  are both bitonic and each element of  $s_1$  is less than every element in  $s_2$ .

- We can apply the procedure recursively on  $s_1$  and  $s_2$  to get the sorted sequence.

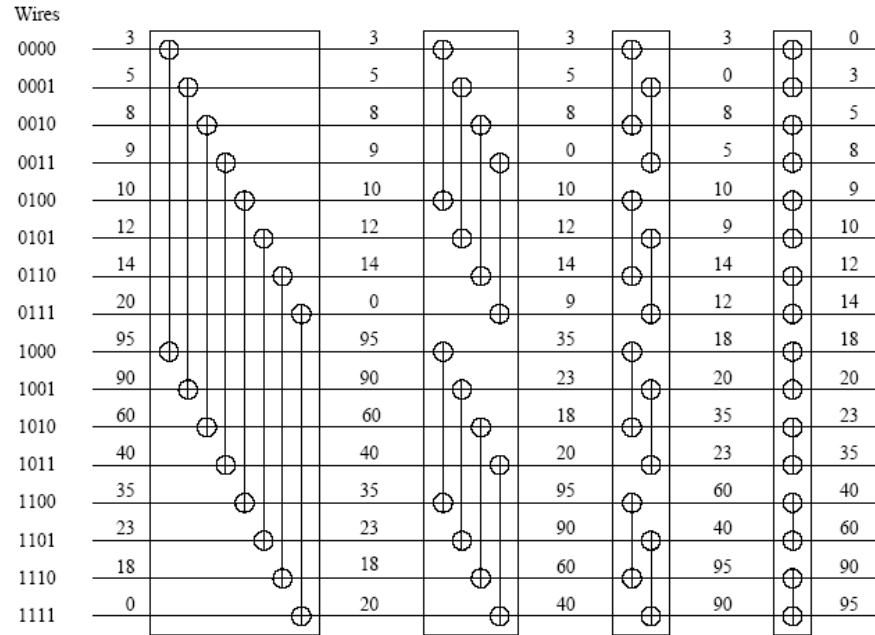
Original sequence	3	5	8	9	10	12	14	20	95	90	60	40	35	23	18	0
1st Split	3	5	8	9	10	12	14	0	95	90	60	40	35	23	18	20
2nd Split	3	5	8	0	10	12	14	9	35	23	18	20	95	90	60	40
3rd Split	3	0	8	5	10	9	14	12	18	20	35	23	60	40	95	90
4th Split	0	3	5	8	9	10	12	14	18	20	23	35	40	60	90	95

Merging a 16-element bitonic sequence through a series of  $\log 16$  bitonic splits.

## Sorting Networks: Bitonic Sort

- We can easily build a sorting network to implement this bitonic merge algorithm.
- Such a network is called a *bitonic merging network*.
- The network contains  $\log n$  columns. Each column contains  $n/2$  comparators and performs one step of the bitonic merge.
- We denote a bitonic merging network with  $n$  inputs by  $\oplus\text{BM}[n]$ .
- Replacing the  $\oplus$  comparators by  $\ominus$  comparators results in a decreasing output sequence; such a network is denoted by  $\ominus\text{BM}[n]$ .

# Sorting Networks: Bitonic Sort




A bitonic merging network for  $n = 16$ . The entire figure represents a  $\oplus\text{BM}[16]$  bitonic merging network. The network takes a bitonic sequence and outputs it in sorted order.

## Sorting Networks: Bitonic Sort

How do we sort an unsorted sequence using a bitonic merge?

We must first build a single bitonic sequence from the given sequence.

- A sequence of length 2 is a bitonic sequence.
- A bitonic sequence of length 4 can be built by sorting the first two elements using  $\oplus\text{BM}[2]$  and next two, using  $\ominus\text{BM}[2]$ .
- This process can be repeated to generate larger bitonic sequences.



A schematic representation of a network that converts an input sequence into a bitonic sequence. In this example,  $\oplus\text{BM}[k]$  and  $\ominus\text{BM}[k]$  denote bitonic merging networks of input size  $k$  that use  $\oplus$  and  $\ominus$  comparators, respectively. The last merging network ( $\oplus\text{BM}[16]$ ) sorts the input. In this example,  $n = 16$ .

# Bitonic Sequence

● Step #1


$a \xrightarrow{\text{red}} = \text{Min}(a,b)$   
 $b \xleftarrow{\text{red}} = \text{Max}(a,b)$

$a \xleftarrow{\text{red}} = \text{Max}(a,b)$   
 $b \xrightarrow{\text{red}} = \text{Min}(a,b)$




# Bitonic Sequence

● Step #2





# Bitonic Sequence

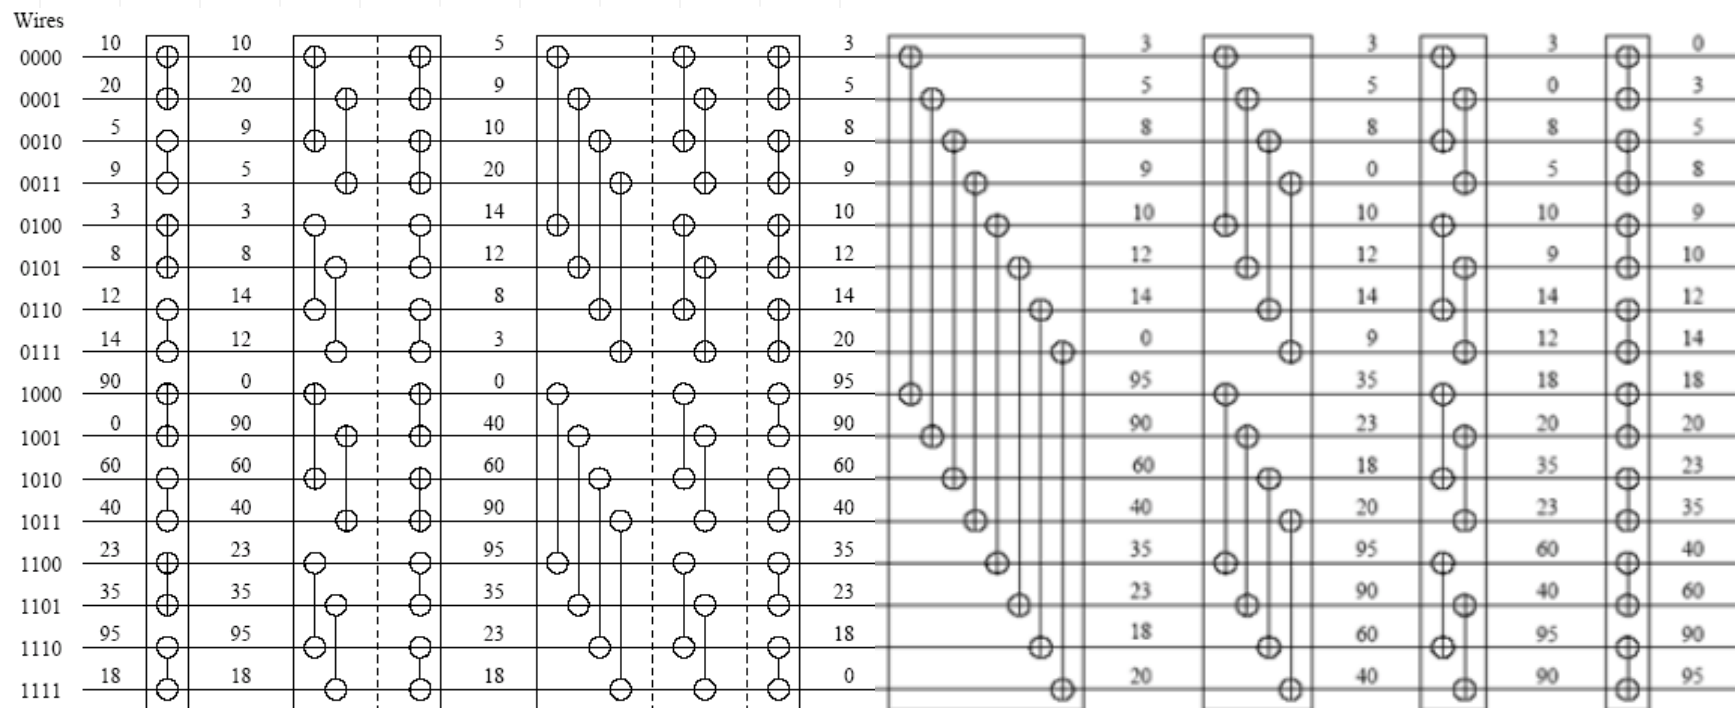


Wires

0000	10
0001	20
0010	5
0011	9
0100	3
0101	8
0110	12
0111	14
1000	90
1001	0
1010	60
1011	40
1100	23
1101	35
1110	95
1111	18



# Sorting Networks: Bitonic Sort



# Bitonic Sort

- The sequence of comparisons is data-independent makes it one of the fastest and suitable parallel sorting algorithms.
- In the first step it makes the arbitrary sequence into a bitonic sequence.
- In the second step the bitonic sequence is sorted.



# Conclusion

- Sorting is one of the most fundamental problems in computer science, as it is used in most software applications.
- This presentation shows different parallel sorting algorithms on GPUs.
- The performance is affected mainly by : nature of algorithm and hardware platforms.